
Bachelor Thesis Final Presentation

Exploring Fuzzy Tuning Technique for Molecular Dynamics Simulations in AutoPas

Manuel Lerchner
manuel.lerchner@tum.de

Advisors:
Manish Kumar Mishra, M.Sc.
Samuel James Newcome, M.Sc.

Table of Contents

1 Introduction

2 Implementation

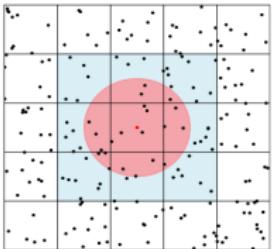
3 Benchmarks

4 Conclusion

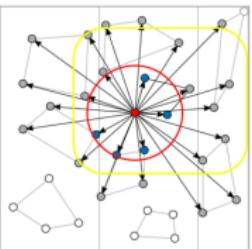
AutoPas

What is AutoPas?

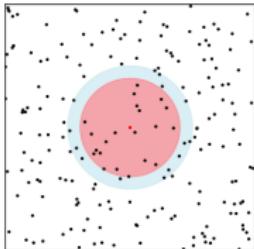
- Library for arbitrary N-body simulations
- Optimal performance by switching implementations
 - **Container:** Finding neighboring particles
 - **Traversal:** Parallel force calculations
 - **Data Layout:** Memory access optimization
 - **Newton 3:** Force calculation optimization



Linked Cells



Verlet Cluster Lists



Verlet Lists

Simpler Memory Access

Lower Memory Overhead

Fewer redundant calculations

[Newcome et al., 2024]

Structure of AutoPas

- Three main areas:
 - User Application
 - Algorithm Library
 - Tuning Strategies
- Algorithm Library:
 - Huge Search Space¹
- Tuning Strategies:
 - Full Search
 - Random Search
 - Predictive Tuning
 - Bayesian Search
 - Rule Based Tuning

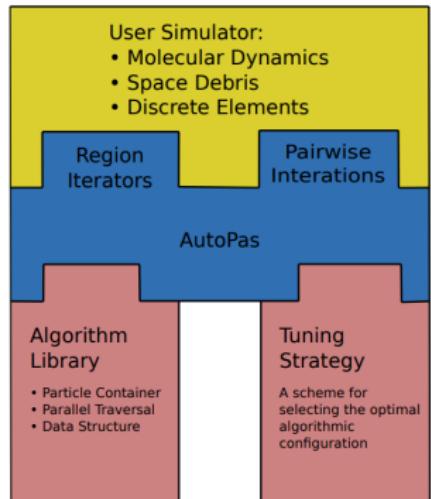
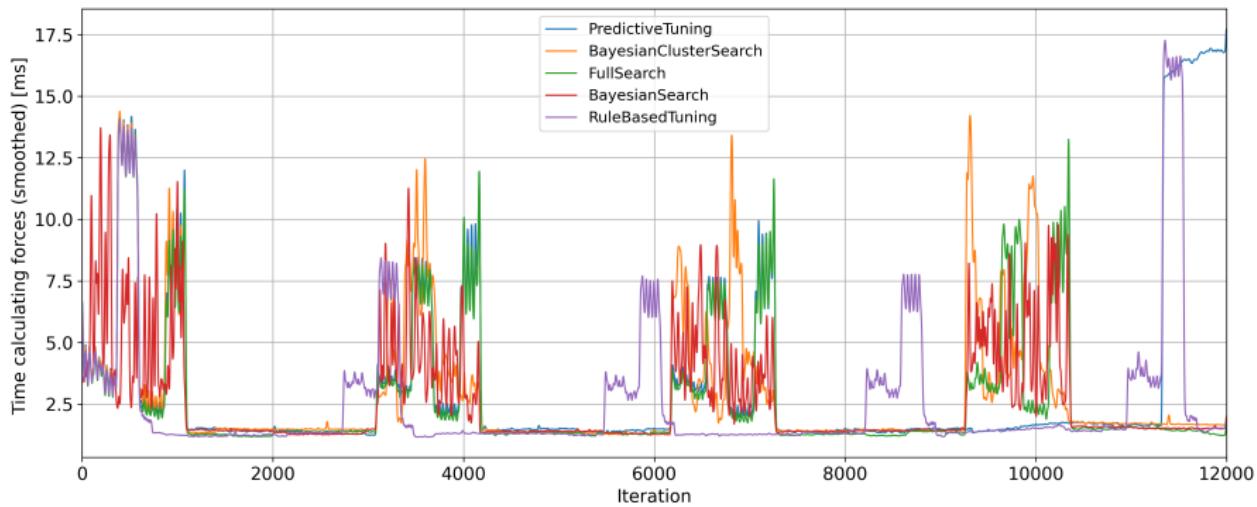


Figure: [Newcome et al., 2023]

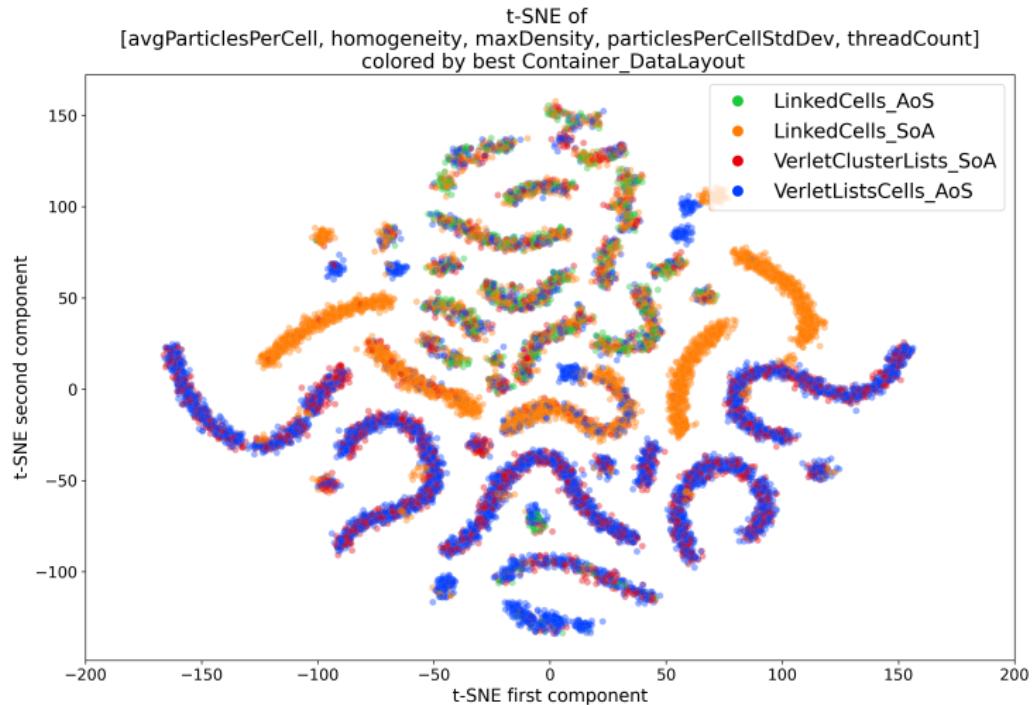
¹Container × Traversal × Data Layout × Newton 3 × Load Estimator × Cell Size Factor

Auto-Tuning

- Tuning Phase → Simulation Phase → Repeat
- Potential Tuning Overhead



Idea: Tuning based on Simulation State



Fuzzy Logic System

- (Fuzzy) Rule-based system
- Human-like reasoning to model systems $f : \mathbb{R}^n \rightarrow \mathbb{R}$
- Linguistic terms instead of numerical values
 - What is *hot*? What is *cold*?
 - Allow smooth transitions between terms

Example (Heater Control)

Input: temperature (e.g. 20°C), humidity (e.g. 60%)

Output: heater power (e.g. 50%)

Rules:

IF temp is <i>cold</i>	AND	humidity is <i>dry</i>	THEN power is <i>high</i>
IF temp is <i>hot</i>	OR	humidity is <i>wet</i>	THEN power is <i>low</i>
IF temp is <i>warm</i>			THEN power is <i>medium</i>

Table of Contents

1 Introduction

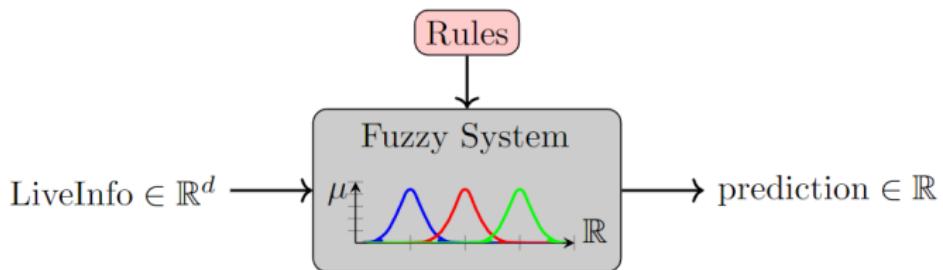
2 Implementation

3 Benchmarks

4 Conclusion

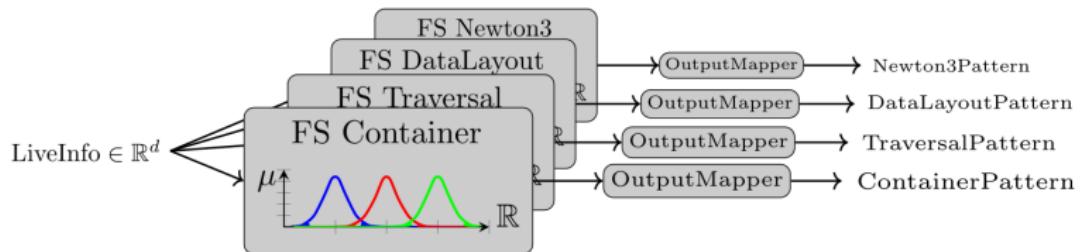
Challenge 1: Numerical Output

- Fuzzy Tuning not directly applicable to AutoPas
 - Tuning is an *algorithm selection* problem
 - Fuzzy tuning provides functions $f : \mathbb{R}^n \rightarrow \mathbb{R}$
 - Big Question: How to interpret the numerical output?



Approach 1: Component Tuning

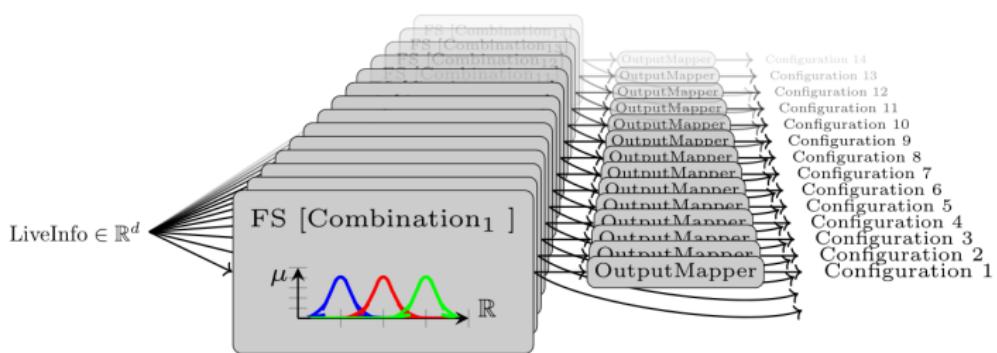
- Explicit Fuzzy System for each tunable parameter
 - Continuous representation of the parameter values



- Example: **IF** avgParticlesPerCell is *low* **AND** threadCount is *low*
THEN traversal is *vcl_c06*

Approach 2: Suitability Tuning

- A Fuzzy System for **each** possible configuration
- Predict *suitability* of each configuration



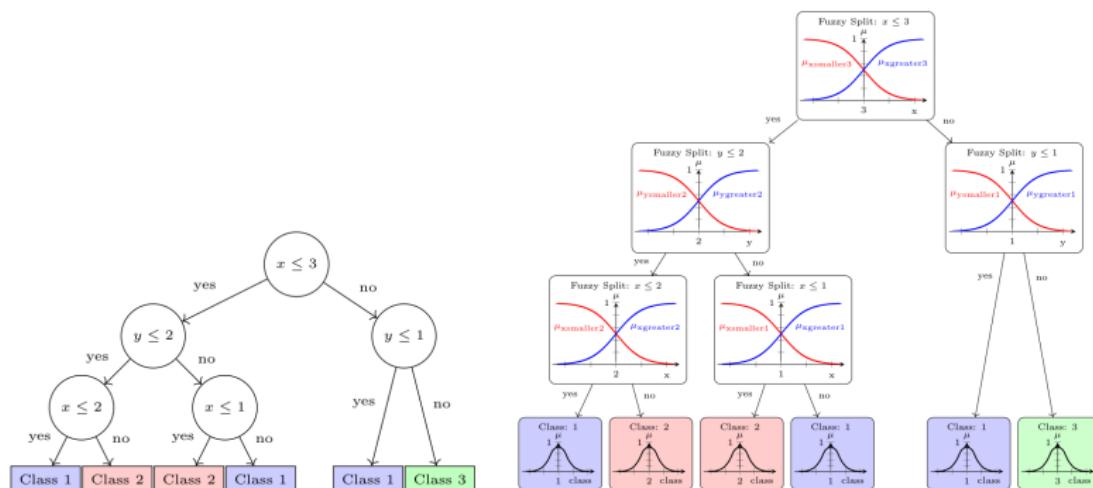
- Example: **IF** `threadCount` is *high* **AND** `avgParticlesPerCell` is *low*
THEN `suitability_LinkedCells_AoS_lc_c18_disabled` is *bad*

Challenge 2: Expert Knowledge

- How to design the Fuzzy System?
 - Creating it manually?
 - ✗ Requires extensive expert knowledge
 - ✗ Difficult to formalize non-trivial knowledge
 - Extracting rules from data?
 - ✓ No prior expert knowledge required
 - ✓ Semi-automated process

Data-Driven Rule Extraction

- Machine learning to generate rules [Crockett et al., 2006]
 - Train Decision Tree on the data
 - Decision Tree → Fuzzy Decision Tree
 - Fuzzy Decision Tree → Fuzzy Rules



Fuzzy Decision Trees → Fuzzy Rules

- Unnest the fuzzy decision tree

Rule	Antecedent	Consequent
1	$x \text{ is smaller3} \wedge y \text{ is smaller2} \wedge x \text{ is smaller2}$	class is 1
2	$x \text{ is smaller3} \wedge y \text{ is smaller2} \wedge x \text{ is greater2}$	class is 2
3	$x \text{ is smaller3} \wedge y \text{ is greater2} \wedge x \text{ is smaller1}$	class is 2
4	$x \text{ is smaller3} \wedge y \text{ is greater2} \wedge x \text{ is greater1}$	class is 1
5	$x \text{ is greater3} \wedge y \text{ is smaller1}$	class is 1
6	$x \text{ is greater3} \wedge y \text{ is greater1}$	class is 3

Fuzzy Rule Extraction for `md_flexible`

- Collect dataset of `md_flexible` simulations
- For each tuning phase i , store:
 - `LiveInfoData`: `maxDensity`, `homogeneity`, `threadCount`, ...
 - `TuningData`: `Container`, `Traversal`, `Newton3`, ..., `Time`
- Introduce *relative speed* metric
 - Absolute time is not meaningful

$$\text{relative speed}_{\text{config}}^{(i)} = \frac{t_{\text{best}}^{(i)}}{t_{\text{config}}^{(i)}}$$

Resulting Dataset

- Contains expected *relative speed* based on:
 - Simulation state
 - Configuration

ParticlesPerCell			Miscellaneous			Configuration			
avg	max	stddev	homogeneity	max-density	threads	Container DataLayout	Traversal	Newton3	Relative Speed
0.905	23	0.0129	0.0354	0.531	1	LinkedCells_AoS	lc_sliced	enabled	0.450641
2.201	13	0.0144	0.0861	0.627	24	VerletListsCells_AoS	vlc_sliced	disabled	0.594117
0.905	18	0.0136	0.0431	0.319	4	LinkedCells_AoS	lc_sliced_c02	enabled	0.454632
:	:	:	:	:	:	:	:	:	:

Fuzzy Rule Extraction

Antecedent				Consequent
avgParticlesPC	homogeneity	particlesPCStdDev	threadCount	Traversal
lower than 1.553	higher than 0.047	lower than 0.023	higher than 2.5	"lc_sliced, vlc_c18, lc_sliced_c02"
	lower than 0.037	lower than 0.023	lower than 26.0	"vlc_c06, vlc_c18, vlc_sliced_c02"
:	:	:	:	:

Example: Component Tuning Rules

Antecedent				Consequent
avgParticlesPC	homogeneity	particlesPCStdDev	threadCount	Suitability
	lower than 0.084	higher than 0.029	higher than 26.0	LinkedCells_AoS lc_c01_disabled "medium"
	higher than 0.084	higher than 0.029	higher than 26.0	"bad"
		higher than 0.02	lower than 2.5	"poor"
:	:	:	:	:

Example: Suitability Tuning Rules

Table of Contents

1 Introduction

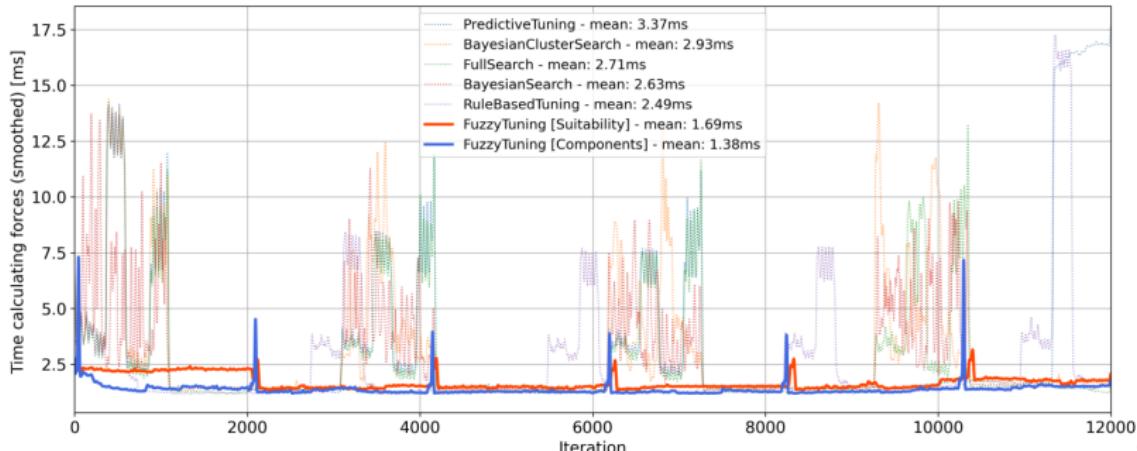
2 Implementation

3 Benchmarks

4 Conclusion

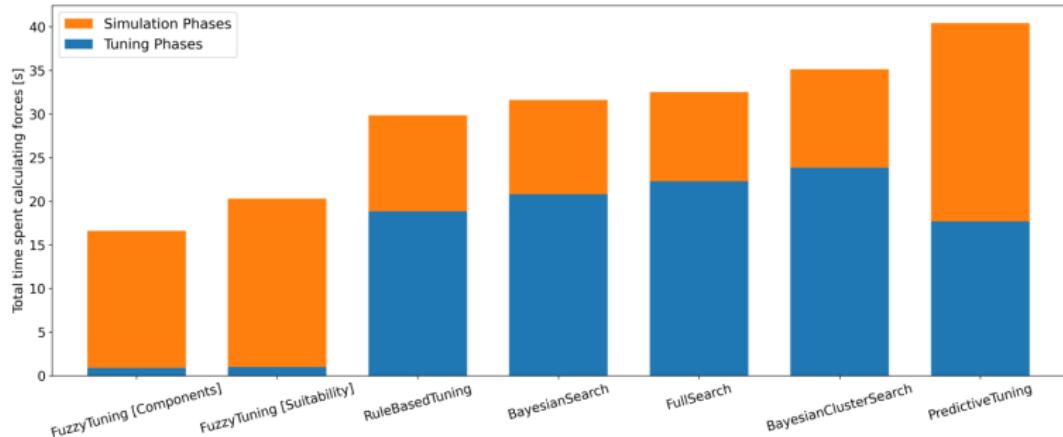
Benchmark 1: Exploding Liquid

- Fuzzy Tuning Approaches:
 - Very short tuning phases
 - Selected configurations perform well
 - Winning configurations are (mostly) equivalent



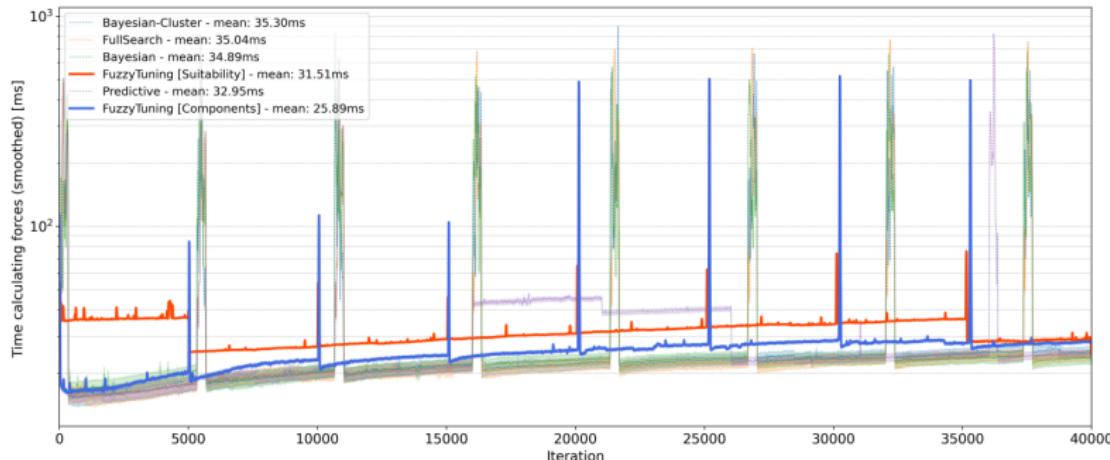
Total Time: Exploding Liquid

- Fuzzy Tuning is good because:
 - Few evaluated configurations
 - Good evaluated configurations
- Massive reduction in time spent tuning



Benchmark 2: Spinodal Decomposition (MPI)

- Component approach performs well
- Suitability approach misses the optimal configuration
 - However: no tuning overhead
 - Can this make up for the suboptimal configurations?



Total Time: Spinodal Decomposition

- Component tuning wins again
- Suitability tuning:
 - Fast tuning phases compensate for suboptimal configurations!
 - Improvement: Encourage longer tuning phases
 - Increase suitability threshold (Top 10% → Top 30%)

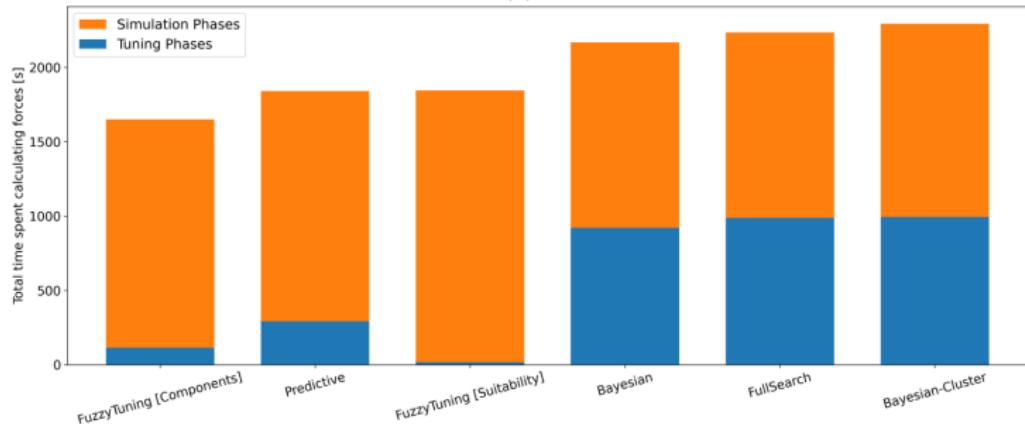


Table of Contents

1 Introduction

2 Implementation

3 Benchmarks

4 Conclusion

Conclusion

- Fuzzy Tuning is very promising for AutoPas
 - Massive reduction in tuning overhead
 - Reduction of total simulation time (up to factor 1.96x)
 - Among the best tuning strategies (but very complex)
- Ongoing Challenge:
 - Rule generation
 - ✗ Requires a lot of data (1.1GB)
 - ✗ Unappealing for users
 - ✗ No universal solution (showed some generalization)

Conclusion

- Fuzzy Tuning is very promising for AutoPas
 - Massive reduction in tuning overhead
 - Reduction of total simulation time (up to factor 1.96x)
 - Among the best tuning strategies (but very complex)
- Ongoing Challenge:
 - Rule generation
 - ✗ Requires a lot of data (1.1GB)
 - ✗ Unappealing for users
 - ✗ No universal solution (showed some generalization)

Improvement Ideas

- Dynamic Rule Generation
 - Update expert knowledge on the fly
 - Automatically adapt to new scenarios
- Improvements to the Tuning Process
 - Investigate *early stopping* mechanism
 - Dismiss slow configurations early
 - Solve tuning overhead once and for all?
- Simplification of the Model
 - Ensemble of Decision Trees instead of Fuzzy Systems
 - Maybe comparable results?

Thank you for your attention!

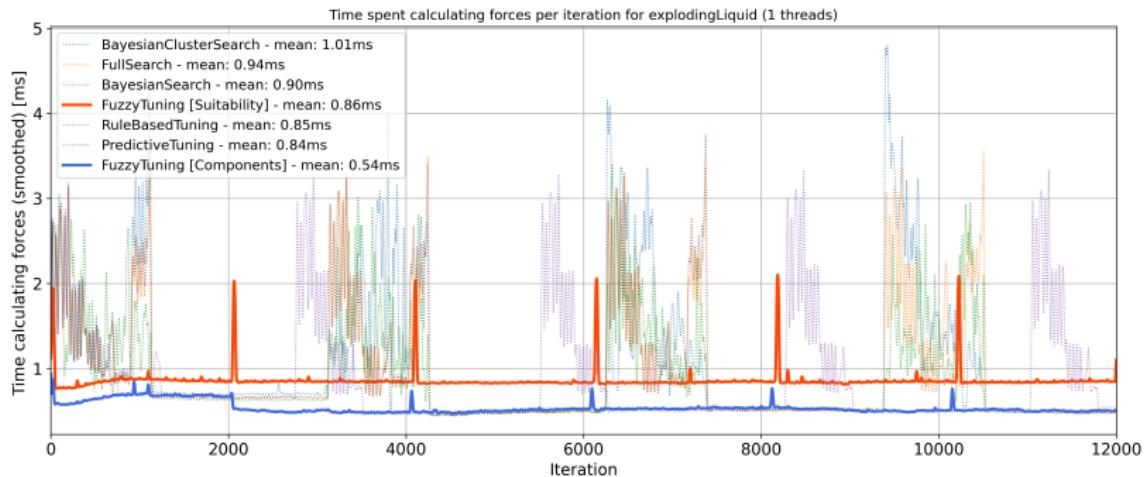
Questions?

References I

-  Crockett, K., Bandar, Z., Mclean, D., and OShea, J. (2006).
On constructing a fuzzy inference framework using crisp decision trees.
Fuzzy Sets and Systems, 157(21):2809–2832.
-  Newcome, S. J., Gratl, F. A., Muehlhaeusser, M., Neumann, P., and Bungartz, H.-J. (2024).
Autopas: Dynamic algorithm selection in molecular dynamics for optimal time and energy.
In *SIAM Conference on Parallel Processing (PP24)*. SIAM.
-  Newcome, S. J., Gratl, F. A., Neumann, P., and Bungartz, H.-J. (2023).
Towards the smarter tuning of molecular dynamics simulations.
Amsterdam, The Netherlands.

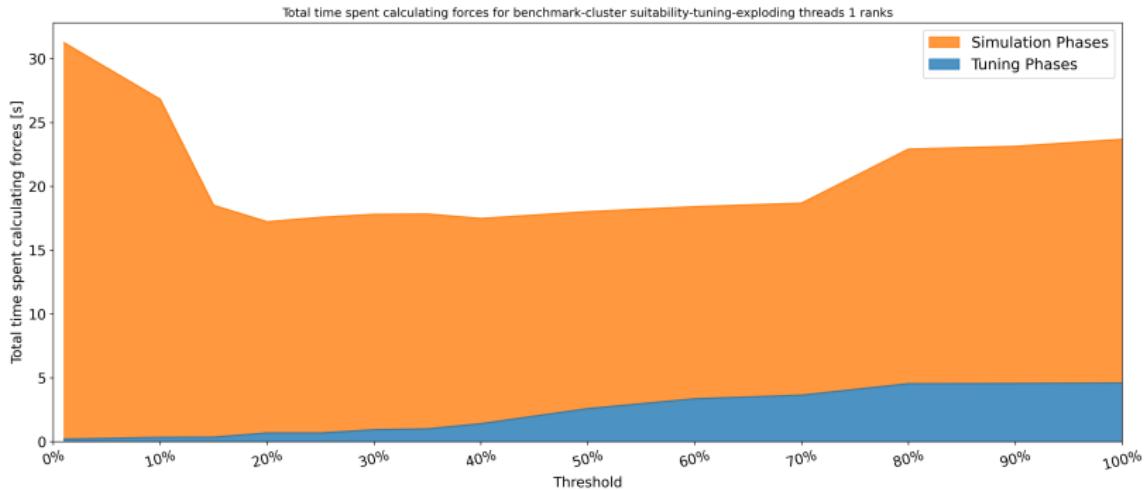
Backup: What happens on new scenarios?

- Exploding Liquid: bigger $\varepsilon \rightarrow$ more spacing
- Component Tuning is winning again

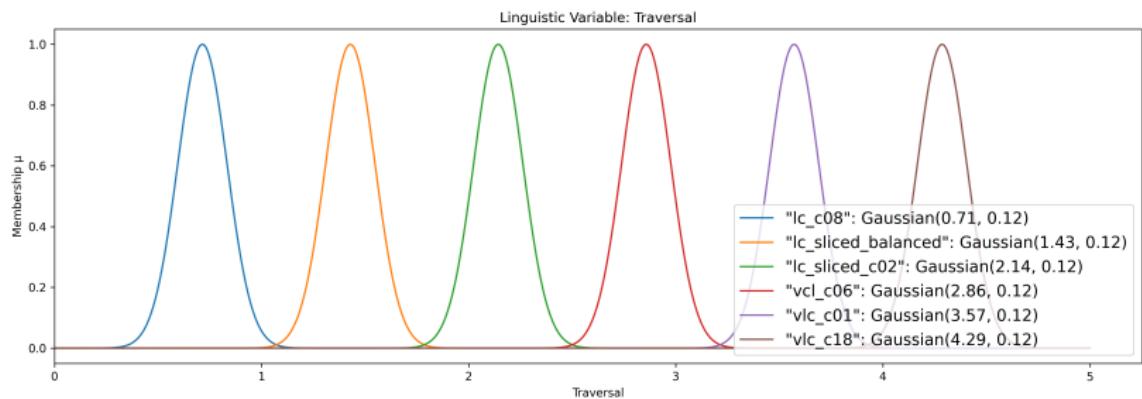


Backup: Optimal Suitability Threshold

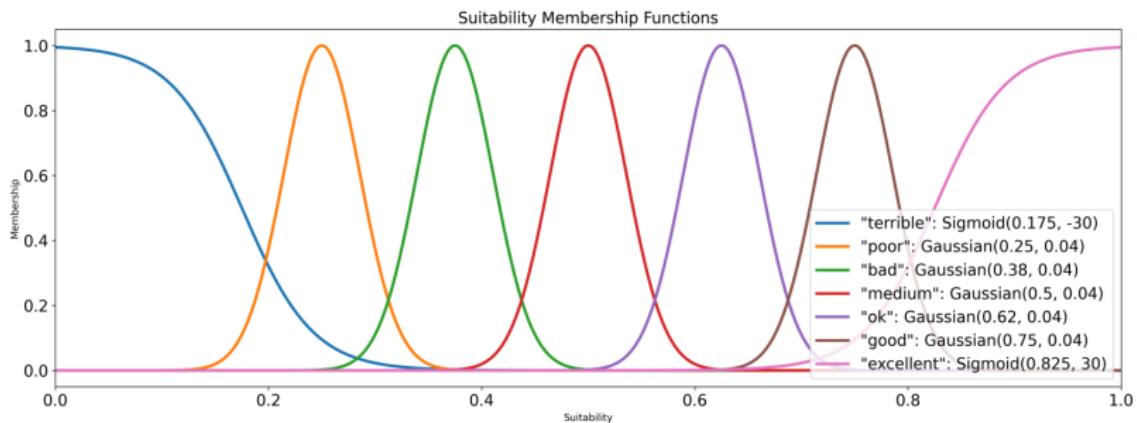
- Goal: find optimal hyperparameters for suitability-approach



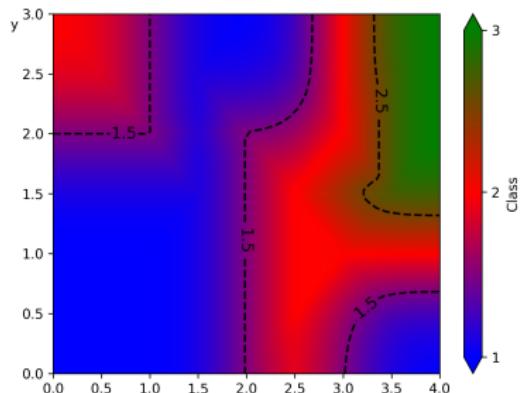
Backup: Linguistic Variables for Component Tuning



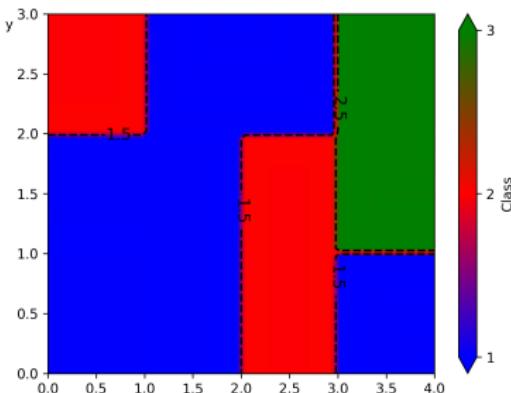
Backup: Linguistic Variables for Suitability Tuning



Backup: Decision Surfaces



(a) Center of Gravity - Defuzzification

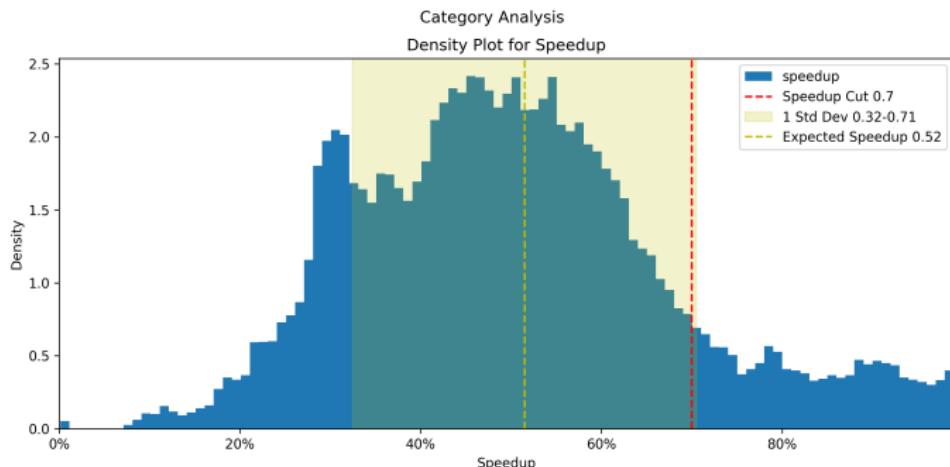


(b) Mean of Maxima - Defuzzification

Figure: Decision Surfaces of Fuzzy System

Backup: Speedup Distribution

- How good is the average configuration?



Backup: Quality of Predictions

- How well do tuning strategies suggest good configurations?

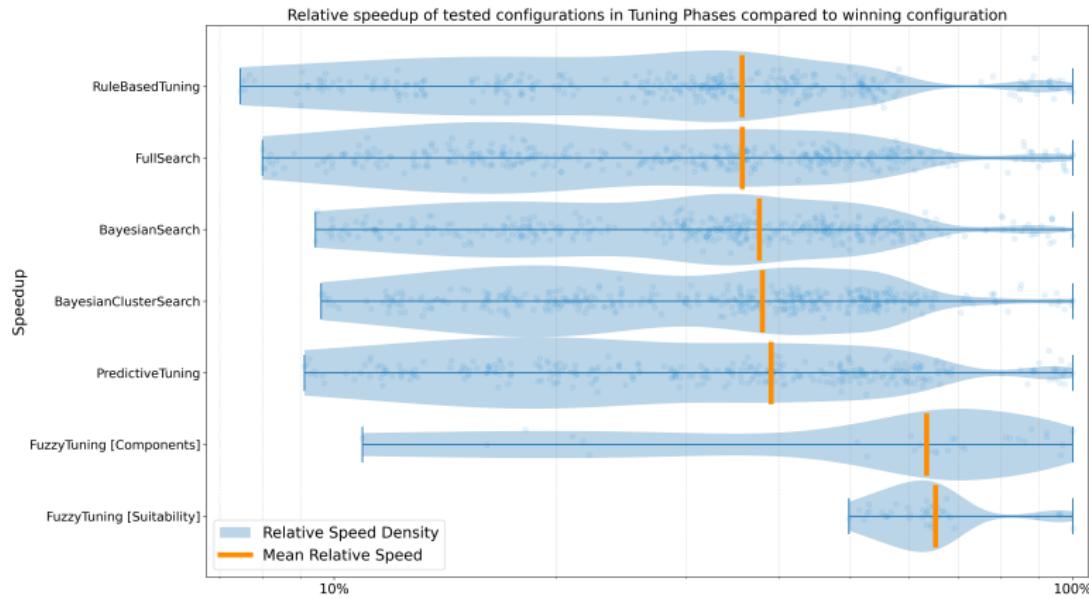


Figure: Source: MathWorks - Fuzzy Inference Process

