
Bachelor Thesis Final Presentation

Exploring Fuzzy Tuning Technique for Molecular Dynamics Simulations in AutoPas

Manuel Lerchner
manuel.lerchner@tum.de

Advisors:
Manish Kumar Mishra, M.Sc. (hons)
Samuel James Newcome, M.Sc.

Table of Contents

- 1 What is AutoPas?
- 2 Mathematics of Fuzzy Logic
- 3 Fuzzy Tuning Strategy for AutoPas
- 4 Approaches for Fuzzy Tuning in AutoPas
 - Component Tuning
 - Suitability Tuning
- 5 Data-Driven Rule Extraction Process
- 6 Fuzzy Rule Extraction for ~~md~~flexible
- 7 Comparison and Evaluation
 - Exploding Liquid
 - Spinodal Decomposition
- 8 Future Work
- 9 Conclusion

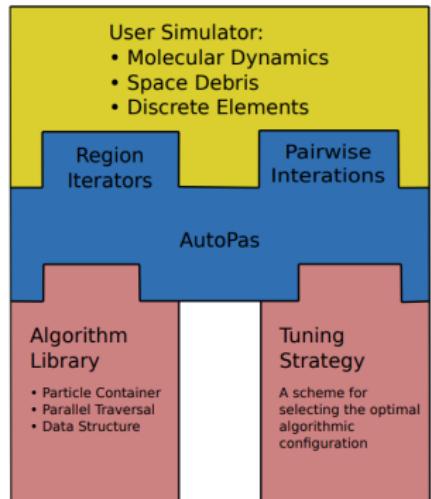


What is AutoPas?

- Library for optimal node-level performance in N-body simulations
- Many different implementations for the N-body problem
- AutoTuning: Automatically switch between implementations
 - **Container:** How to find neighboring particles?
 - **Traversal:** How to handle multi-threading?
 - **Data Layout:** How to store particles in memory?
 - **Newton 3:** Can we exploit Newton's 3rd law?
 - ...
- Official Example applications:
 - `md_flexible` (Molecular Dynamics)
 - `sph` (Smoothed Particle Hydrodynamics)

Structure of AutoPas

- Three main components:
 - User Application
 - Algorithm Library
 - Tuning Strategies
- Algorithm Library:
 - Huge Search Space¹
- Tuning Strategies:
 - Full Search
 - Random Search
 - Predictive Tuning
 - Bayesian Search
 - Rule Based Tuning

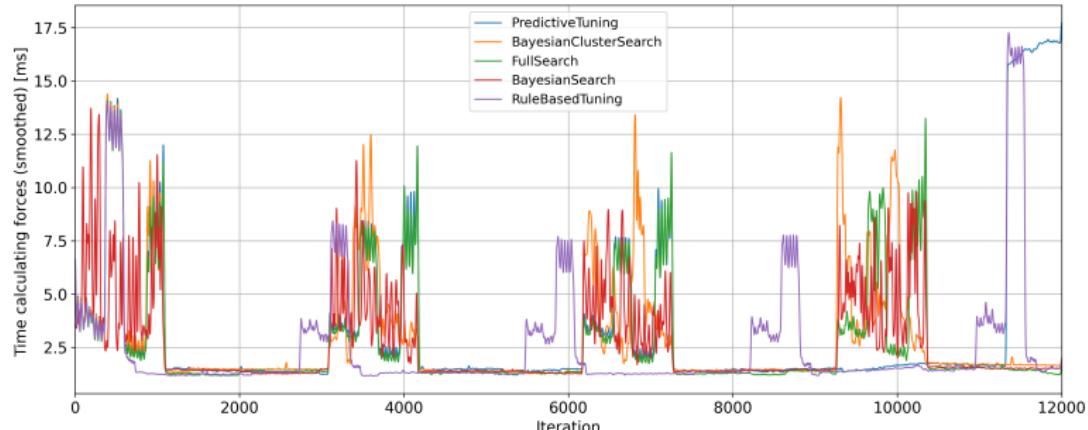


Source: [3]

¹Container × Traversal × Data Layout × Newton 3 × Load Estimator × Cell Size Factor

Auto-Tuning

- Tuning Phase: Find the best configuration
 - Tuning Strategies select configurations to evaluate
 - Expensive, Time consuming
- Simulation Phase: Use the best configuration



Fuzzy Logic Systems

- Use human-like reasoning to model complex systems
- Example: Heater Control
 - Input: temperature (e.g. 20°C), humidity (e.g. 50%)
 - Output: heater power (e.g. 50%)
 - Rules:

IF temp is <i>cold</i>	AND	humidity is dry	THEN power is <i>high</i>
IF temp is <i>hot</i>	OR	humidity is wet	THEN power is <i>low</i>
IF temp is <i>warm</i>		THEN power is <i>medium</i>	
- Very easy to understand and interpret
- Can handle uncertainty and imprecise information
- Complexity is abstracted away in the linguistic terms (e.g. *cold*, *warm*, *hot*)
- Fuzzy System $f : \mathbb{R}^n \rightarrow \mathbb{R}$

Mathematical Foundations

- Consider the Fuzzy Rule:

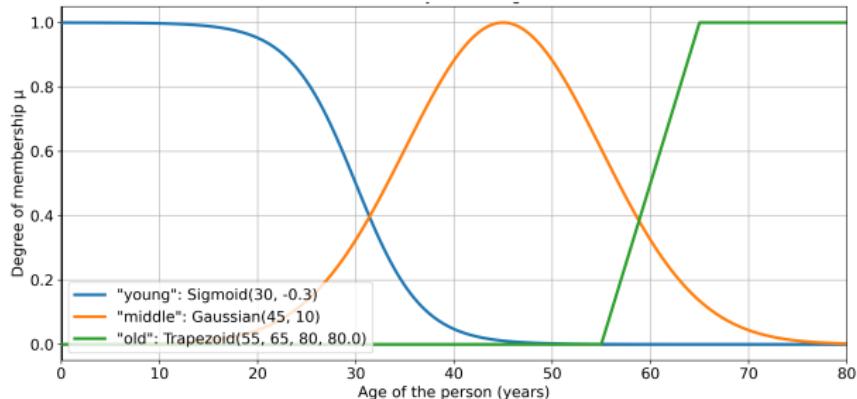
The diagram illustrates a Fuzzy Rule structure. It consists of two main horizontal lines. The upper line represents the rule's components: 'IF' followed by a set of conditions enclosed in parentheses, and 'THEN' followed by a consequent. The conditions are: '(temp is *cold* AND hum is *dry*)'. The consequent is 'power is *high*'. Below the first part of the rule ('IF ...') are labels 'Ling. Var.' and 'Ling. Term' under each of the two terms (*cold* and *dry*). Below the second part ('THEN ...') are also 'Ling. Var.' and 'Ling. Term' labels. The middle section between the two parts is labeled 'Antecedent'. The lower section is labeled 'Consequent'. The entire rule is bracketed at the bottom as a 'Fuzzy Rule'.

- Fuzzy Logic Systems consist of:

- Linguistic Terms / Fuzzy Sets (e.g. *cold*, *warm*, *hot*)
 - Linguistic Variables (e.g. temperature, humidity, power)
 - Fuzzy Logic Operators (e.g. **AND**, **OR**, **NOT**)
 - Fuzzy Rules (e.g. **IF** antecedent **THEN** consequent)

Fuzzy Sets

- Fuzzy Sets are generalizations of classical sets
 - Classical Sets: binary membership function $\in_A: A \rightarrow \{\text{false}, \text{true}\}$
- Fuzzy Sets are defined by:
 - Underlying Crisp Set X (e.g. $\text{age} \subseteq \mathbb{R}$)
 - **Continuous** membership function $\mu_{\tilde{A}}: X \rightarrow [0, 1]$
- Allow for uncertainty. When is a person young?



Linguistic Variables

- Linguistic Variables can take on linguistic terms / fuzzy sets
 - E.g. *age* can take *young*, *middle-aged* or *old*
- Instead of using crisp values (35 years), we use a combination of linguistic terms to describe the age (Fuzzyification):

$$35 \text{ years} \implies \begin{cases} 20\% \text{ young} \\ 60\% \text{ middle-aged} \\ 0\% \text{ old} \end{cases}$$

- This allows for non-numerical reasoning:
 - E.g. **IF** age is *young* **THEN** fitness is *high*
 - Use abstract concepts instead of crisp values
- Each linguistic term represents a certain *collection* of values

Fuzzy Logic Operators

- Fuzzy Logic Operators are used to modify/combine fuzzy sets
- Extension of boolean logic operators to real numbers
 - $\wedge : \{\text{false}, \text{true}\} \times \{\text{false}, \text{true}\} \rightarrow \{\text{false}, \text{true}\}$
 - **AND** : $[0, 1] \times [0, 1] \rightarrow [0, 1]$
- Extended operators need to maintain the classical semantics
- Typically, Fuzzy Logic Operators are defined as:
 - **AND**: Corresponds to the intersection of fuzzy sets

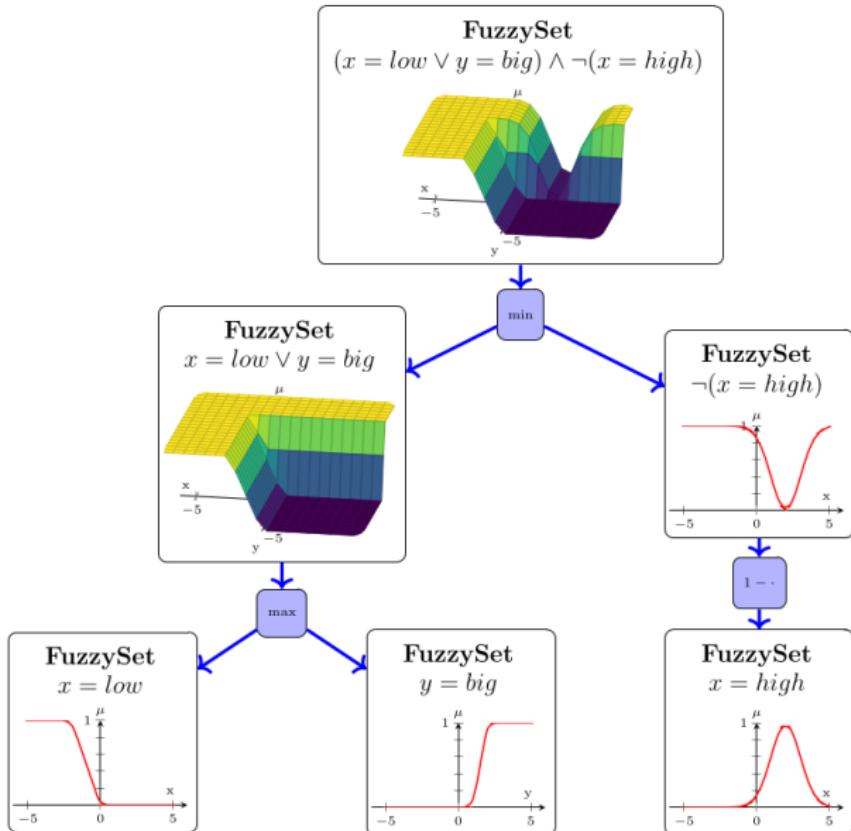
$$\mu_{\tilde{A} \cap \tilde{B}}(x) = \min(\mu_{\tilde{A}}(x), \mu_{\tilde{B}}(x))$$

- **OR**: Corresponds to the union of fuzzy sets

$$\mu_{\tilde{A} \cup \tilde{B}}(x) = \max(\mu_{\tilde{A}}(x), \mu_{\tilde{B}}(x))$$

- **NOT**: Corresponds to the complement of a fuzzy set

$$\mu_{\neg \tilde{A}}(x) = 1 - \mu_{\tilde{A}}(x)$$



Fuzzy Rules

- Each rule is of the form: **IF antecedent THEN consequent**
 - Both antecedent and consequent are fuzzy sets
 - E.g. **IF age is young THEN fitness is high**
- Rules can be interpreted as a logical implication
 - **IF \tilde{A} THEN \tilde{B}** $\iff \tilde{A} \text{ IMPLIES } \tilde{B}$
 - Implication is similar to previous operators (**AND, OR, NOT**)
 - Special form of implication: Mamdani Implication
 - $\tilde{R} = \text{IF } \tilde{A} \text{ THEN } \tilde{B}$
 - $\mu_{\tilde{R}}(x) = \min(\mu_{\tilde{A}}(x), \mu_{\tilde{B}}(x))$
 - *Effect of the rule is limited by the strength of the antecedent*
- A Fuzzy System can consist of multiple rules acting on the same linguistic variable
 - The total *effect* on the output is the combination/union of all individual rule outputs

Defuzzification

- Process of converting arbitrary fuzzy sets to a crisp value
 - Special case: Fuzzy sets resulting from rule application

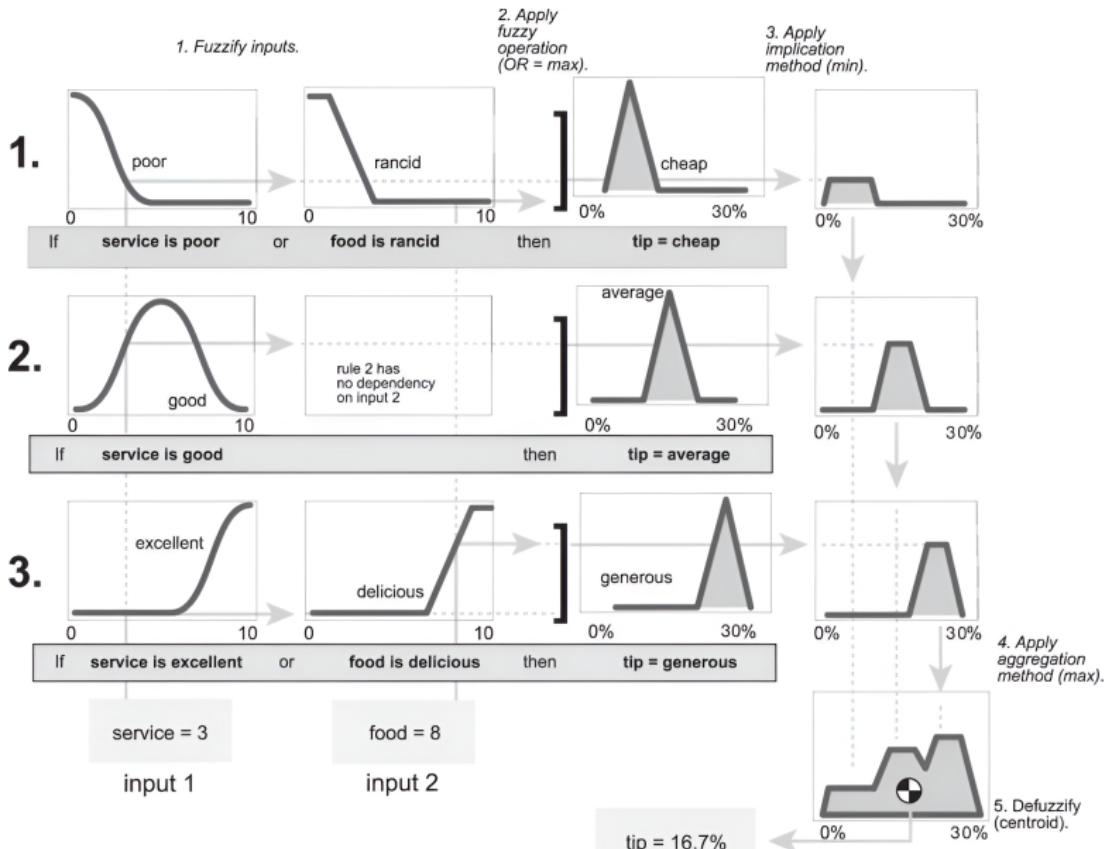
$$\begin{cases} 20\% \text{ young} \\ 60\% \text{ middle-aged} \\ 0\% \text{ old} \end{cases} \implies 35 \text{ years}$$

- Multiple methods for defuzzification:
 - **Centroid:** Weighted average of the fuzzy set

$$\text{Centroid} = \frac{\int x \cdot \mu_{\tilde{R}}(x) dx}{\int \mu_{\tilde{R}}(x) dx}$$

- **Mean of Maxima:** Average of all maxima of the fuzzy set
- Core idea: Represent aspects of the fuzzy set with a crisp value
 - E.g. Weighted average of all possible values (Centroid)
 - E.g. Most likely value (Mean of Maxima)

Source: MathWorks - Fuzzy Inference Process



- TODO: Maybe add decision surfaces here

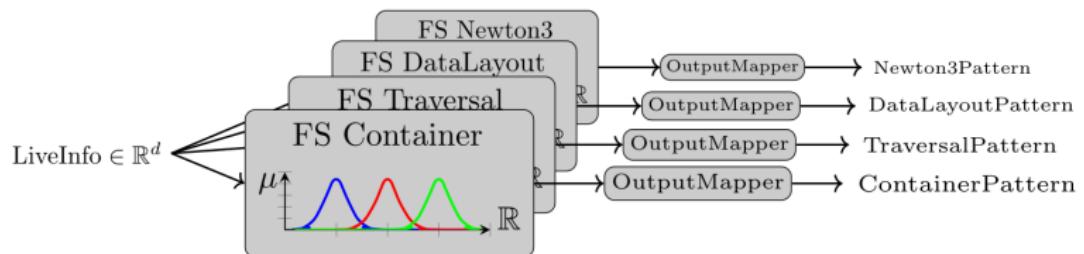
Fuzzy Tuning Strategy

- Main Idea: Use Fuzzy Logic to tune AutoPas
- Make use of LiveInfoData² to perform tuning
- Benefits:
 - Similar to Rule-Based Tuning
 - Potentially more expressive and powerful
 - Still easy to understand and interpret
- Challenges and Questions for AutoPas:
 - What are the output variables? How to predict Configurations?
 - How to interpret the result? (Fuzzy System : $f : \mathbb{R}^n \rightarrow \mathbb{R}$)
 - How to create the fuzzy rules? Expert knowledge?
 - How to specify the linguistic terms / fuzzy sets?

²Simulation state: avgParticles/Cell, homogeneity, threadCount ...

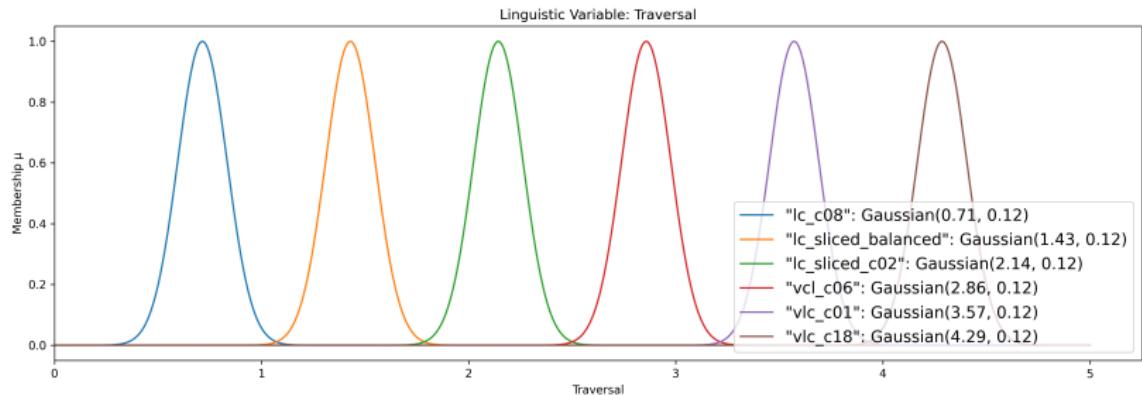
Approach 1: Component Tuning

- Independently predict good values/patterns for each tunable parameter
- Create a Fuzzy System for each tunable parameter
 - Container_DataLayout
 - Traversal
 - Newton 3
- Output Variables: Nominal Representation of parameter values
- Numeric output mapped to the closest value [2]
- Combine the patterns to obtain final configurations



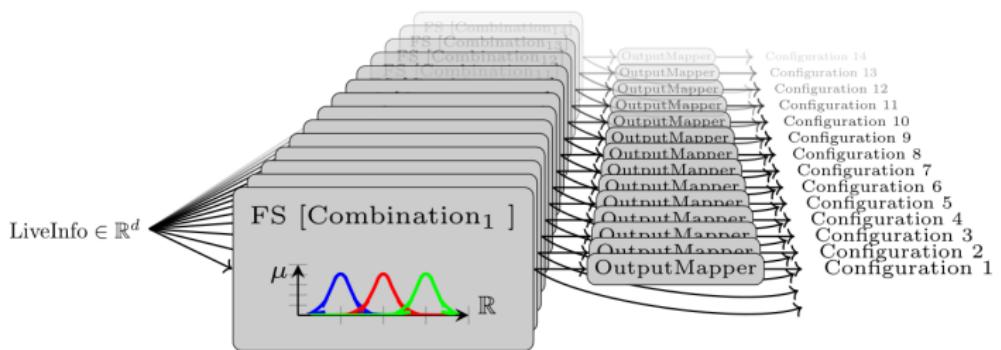
Approach 1: Component Tuning

- Each parameter has a separate Fuzzy System
- Mean of Maxima defuzzification \implies no interpolation
- Example rule:
 - **IF** avgParticlesPerCell is *low* **AND** threadCount is *low*
THEN Traversal is *vcl_c06*
 - Where *low*, *high*, *AoS* are appropriate linguistic terms



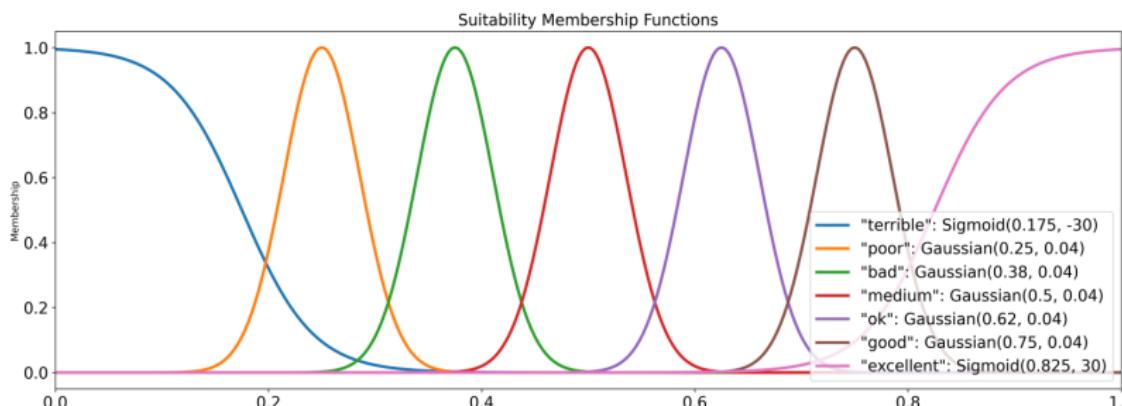
Approach 2: Suitability Tuning

- Predict the suitability of *each* configuration
- Use the suitability values to determine worthwhile configurations
- Create a Fuzzy System for each possible configuration
- Output Variables: Suitability of the configuration
- More complex, but potentially more powerful



Approach 2: Suitability Tuning

- Each configuration has a separate suitability variable
- Center of Gravity defuzzification
- Example rule:
 - **IF** `threadCount` is *high* **AND** `avgParticlesPerCell` is *low*
THEN `suitability_LinkedCells_AoS_lc_c18_disabled` is *bad*
 - Where *high*, *low*, *bad* are appropriate linguistic terms



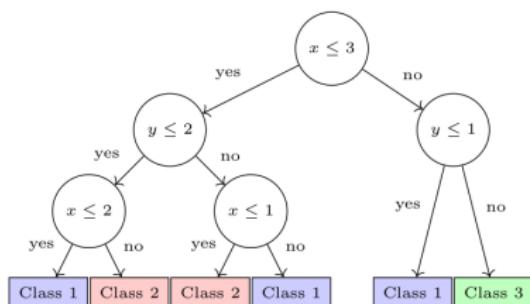
Data-Driven Rule Extraction

- Creating the Fuzzy Rules is hard
 - Expert Knowledge is required
 - Formalization of knowledge is difficult
 - Potentially many rules required
- Use Machine Learning to extract rules from data
- Decision Tree → Fuzzy Decision Tree → Fuzzy Rules [1]
- Does not require expert knowledge
- Human experts can still validate the rules

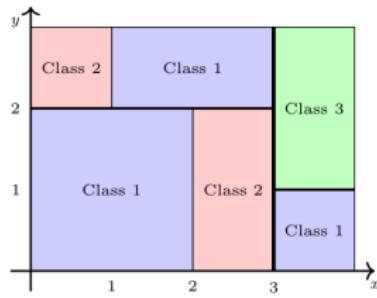
Conversion Process

- Decision Trees

- Try to find the best axis-aligned splits to separate classes
- Results in *if-then-else* rules
- Easy to understand and interpret
- Many libraries available
- Data → Decision Tree (trivial)



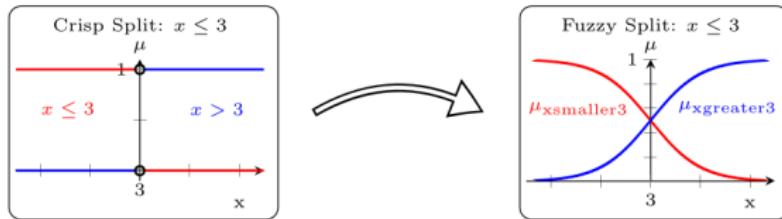
(a) Example decision tree

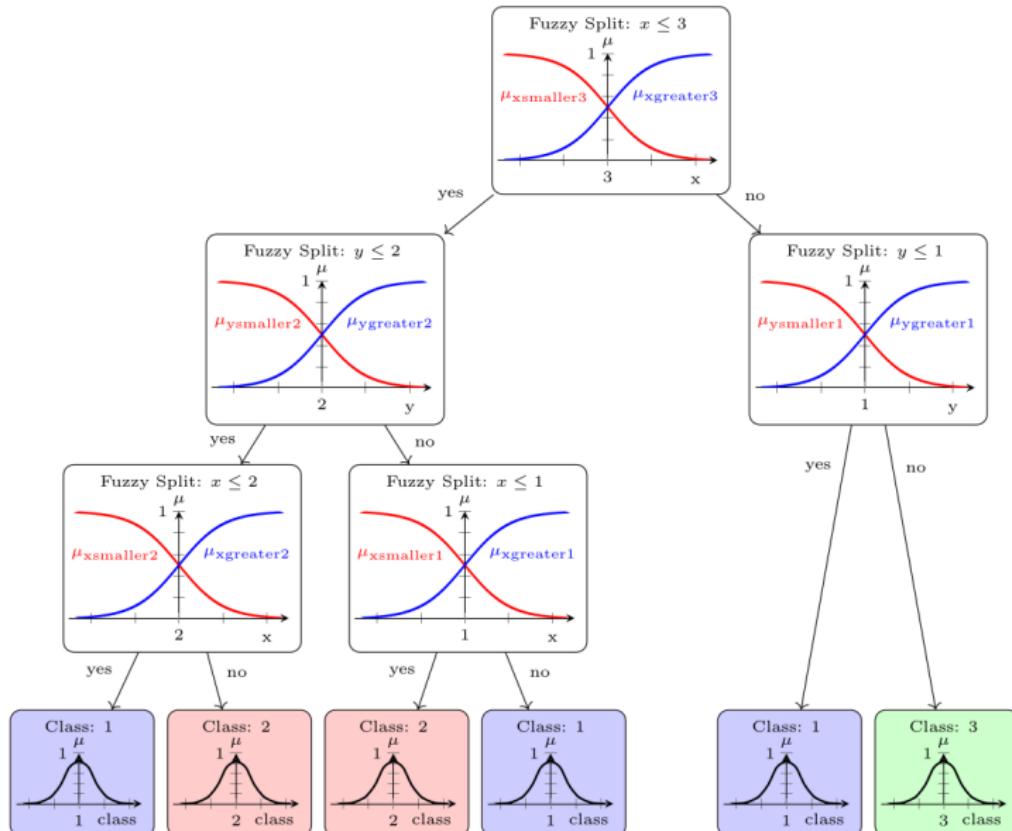


(b) Decision surface over $\mathcal{D} = [0, 4] \times [0, 3]$

Conversion Process

- Fuzzy Decision Trees
 - Convert each (crisp) split into two fuzzy sets
 - Fuzzy sets should maintain the semantics of the split
 - Introduce uncertainty in the split
 - Helpful for generalization, especially if value is close to the split
 - Convert leaf nodes into linguistic terms
- Create linguistic terms for all Leaf Nodes
 - Each class represented by a linguistic term
 - Collect them into a Linguistic Variable





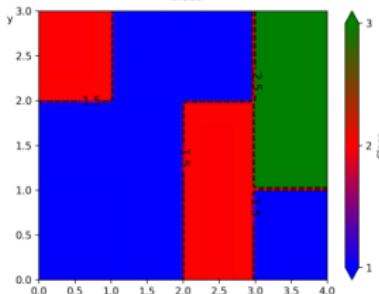
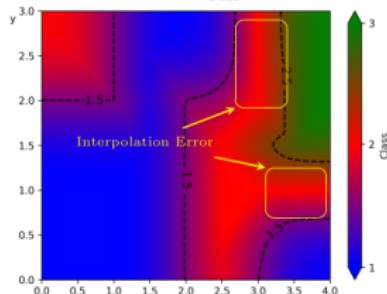
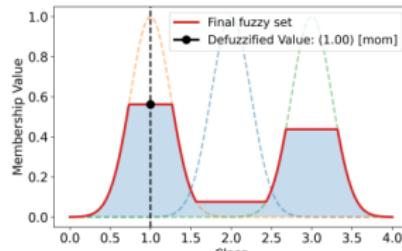
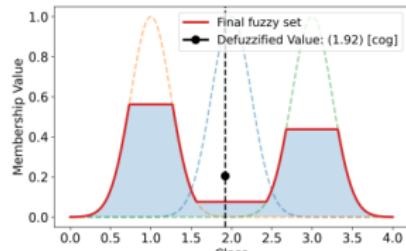
Fuzzy Rule Extraction

- Depth-First traversal of the Fuzzy Decision Tree
 - Each left-right decision corresponds to a condition
 - Leaf node corresponds to the output
- One rule per path from root to leaf

Rule	Antecedent	Consequent
1	$x \text{ is smaller3} \wedge y \text{ is smaller2} \wedge x \text{ is smaller2}$	$class \text{ is } 1$
2	$x \text{ is smaller3} \wedge y \text{ is smaller2} \wedge x \text{ is greater2}$	$class \text{ is } 2$
3	$x \text{ is smaller3} \wedge y \text{ is greater2} \wedge x \text{ is smaller1}$	$class \text{ is } 2$
4	$x \text{ is smaller3} \wedge y \text{ is greater2} \wedge x \text{ is greater1}$	$class \text{ is } 1$
5	$x \text{ is greater3} \wedge y \text{ is smaller1}$	$class \text{ is } 1$
6	$x \text{ is greater3} \wedge y \text{ is greater1}$	$class \text{ is } 3$

Fuzzy Decision Surfaces - CoG vs. MoM

- CoG: Interpolation effect + errors, smooth boundaries
- MoM: Hard boundaries, similar to Decision Trees



Fuzzy Rule Extraction for `md_flexible`

- Collect huge dataset of `md_flexible` simulations
- **LiveInfoData:** `maxDensity`, `homogeneity`, `threadCount`, ...
- **TuningData:** `Container`, `Traversal`, `Newton3`, ..., `Time`
- Introduce notion of *relative speed* for each configuration
 - $t_{best}^{(i)}$: Best configuration time in tuning phase i
 - $t_{config}^{(i)}$: Time of configuration in tuning phase i

$$\text{relative speed}_{config}^{(i)} = \frac{t_{best}^{(i)}}{t_{config}^{(i)}}$$

- Allows for fair comparison between different tuning phases

Resulting Dataset

- Performance of configurations in different environments
- Can be used to extract rules for both approaches
- Goal: Find configurations with high relative speed

ParticlesPerCell			Miscellaneous			Configuration			Relative Speed
avg	max	stddev	homogeneity	max-density	threads	Container DataLayout	Traversal	Newton3	
0.905	23	0.0129	0.0354	0.531	1	LinkedCells_AoS	lc_sliced	enabled	0.450641
2.201	13	0.0144	0.0861	0.627	24	VerletListsCells_AoS	vlc_sliced	disabled	0.594117
0.905	18	0.0136	0.0431	0.319	4	LinkedCells_AoS	lc_sliced_c02	enabled	0.454632
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮

Rulefile Approach 1: Component Tuning

- Naively remove bad configurations (relative speed < 70%)
- Group remaining dataset by tuning-phase
- Aggregate present (*good*) values for each parameter
- Apply Rule Extraction algorithm to each parameter

ParticlesPerCell			Miscellaneous			Aggregated Configuration Terms		
avg	max	stddev	homogeneity	max-density	threads	Container DataLayout	Traversal	Newton3
0.906	15	0.015	0.055	0.297	4	"LinkedCells_SoA, VerletClusterLists_SoA, VerletListsCells_AoS"	"lc_sliced, lc_sliced_balanced, lc_sliced_c02"	"enabled"
0.945	25	0.041	0.084	0.673	24	"LinkedCells_SoA, VerletClusterLists_SoA, VerletListsCells_AoS"	"lc_c04, lc_c08, lc_sliced, lc_sliced_balanced"	"disabled, enabled"

Antecedent

Consequent

avgParticlesPC	homogeneity	particlesPCStdDev	threadCount	Traversal
lower than 1.553	higher than 0.047	lower than 0.023	higher than 2.5	"lc_sliced, vlc_c18, lc_sliced_c02"
	lower than 0.037	lower than 0.023	lower than 26.0	"vlc_c06, vlc_c18, vlc_sliced_c02"
:	:	:	:	:

Rulefile Approach 2: Suitability Tuning

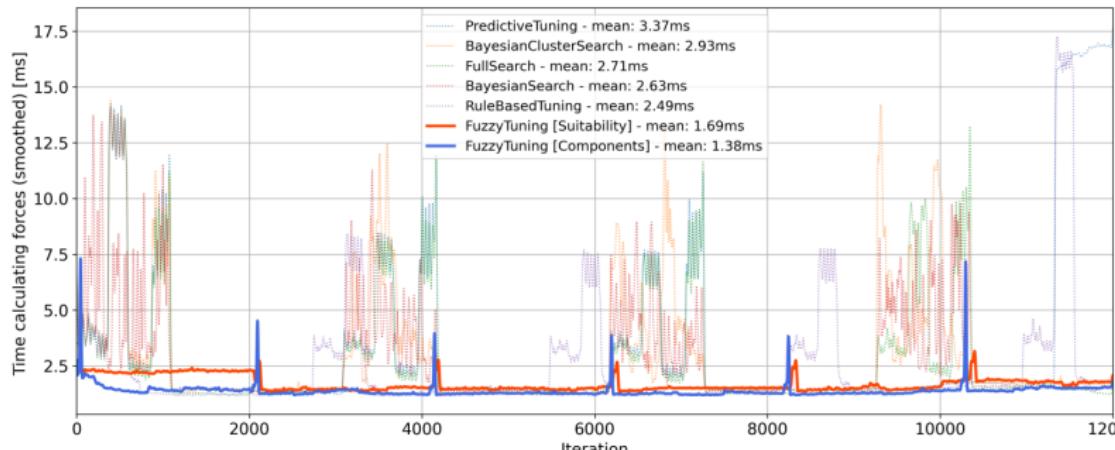
- Assign suitability-classes to each entry in the dataset
- Classes represent ranges of relative speed
- Apply Rule Extraction algorithm for the suitability class

ParticlesPerCell			Miscellaneous			Configuration				
avg	max	stddev	homogeneity	max-density	threads	Container DataLayout	Traversal	Newton3	Relative speed	Suitability
0.905	15	0.012	0.035	0.531	1	LinkedCells_AoS	lc sliced	enabled	0.450	"bad"
0.944	25	0.012	0.083	0.691	28	VerletClusterLists_AoS	vcl_c06	disabled	0.319	"poor"
0.944	20	0.012	0.079	0.041	12	LinkedCell_SoA	vlc_ sliced	enabled	0.989	"excellent"

Antecedent				Consequent
avgParticlesPC		homogeneity	particlesPCStdDev	threadCount
		lower than 0.084	higher than 0.029	higher than 26.0
		higher than 0.084	higher than 0.029	higher than 26.0
			higher than 0.02	lower than 2.5
⋮	⋮	⋮	⋮	⋮

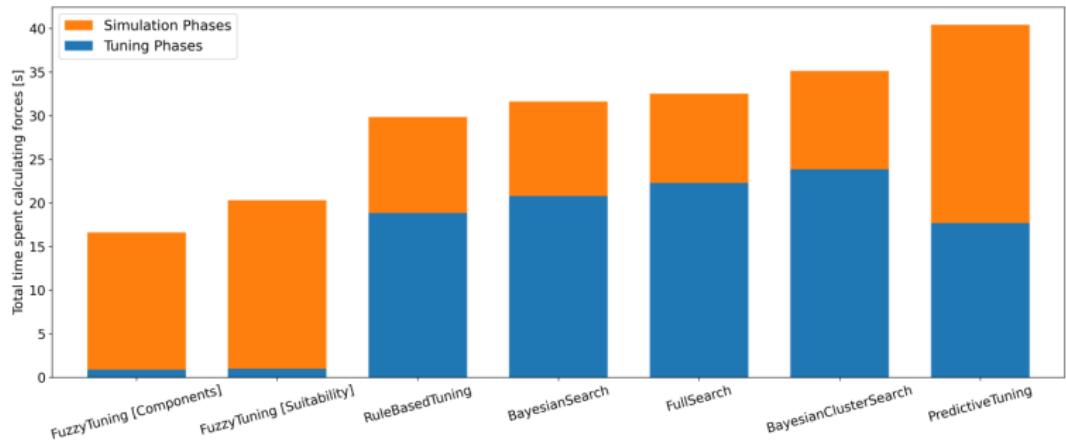
Comparison and Evaluation: Exploding Liquid

- Exploding Liquid Benchmark (Included in Dataset)
 - Both approaches look promising
 - Very short tuning phases (not much noise)
 - Selected configurations perform well (tiny spikes)
 - Winning configurations are (mostly) equivalent



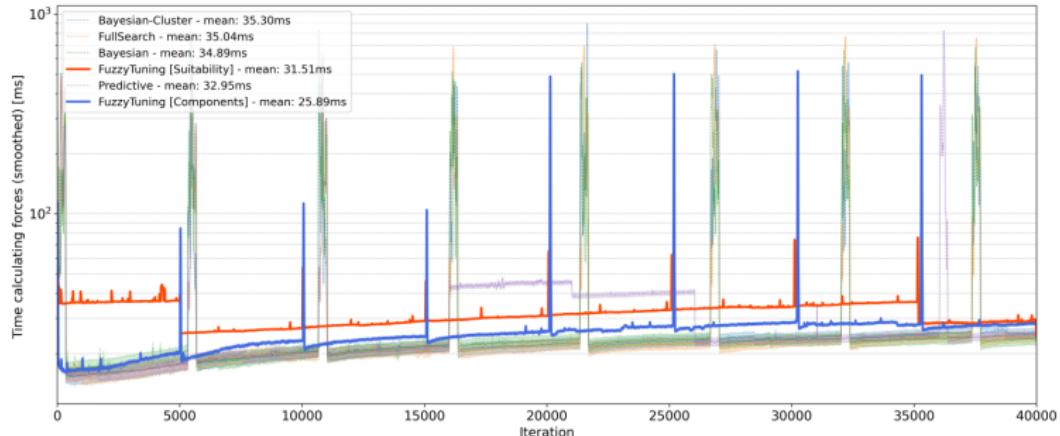
Comparison and Evaluation: Exploding Liquid

- Exploding Liquid Benchmark (Included in Dataset)
 - Fuzzy tuning performs the best by far
 - Mostly due to the very efficient tuning phases
 - Few configurations evaluated
 - Evaluated configurations are expected to perform well
 - Benefits of tuning with tiny overhead



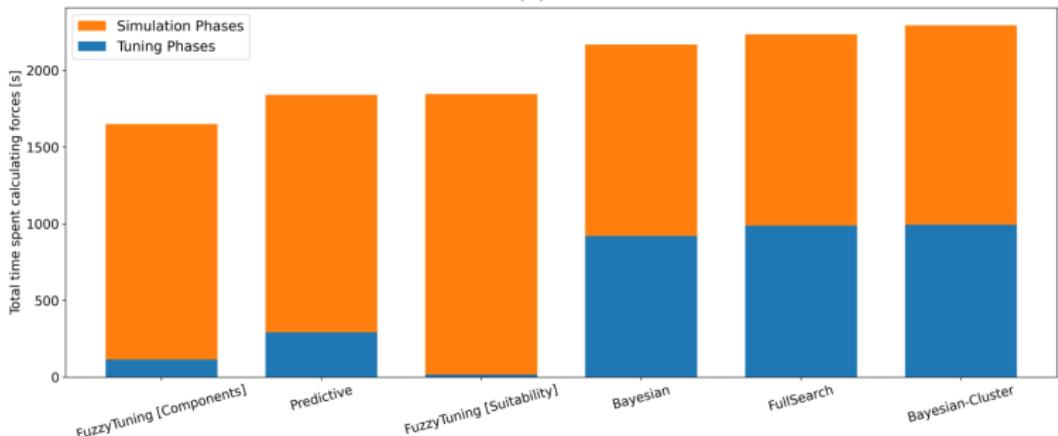
Comparison and Evaluation: Spinodal Decomposition MPI

- Spinodal Decomposition (Included in Dataset)
 - Fuzzy tuning promising again
 - Suitability approach struggles, often misses the best configuration
 - However, both approaches have efficient tuning phases



Comparison and Evaluation: Spinodal Decomposition

- Spinodal Decomposition (Indirectly included in Dataset)
 - Fuzzy tuning performs relatively well
 - Suitability approach still quite good, despite never finding the best configuration
 - Classical tuning strategies perform poorly in tuning phases
 - Fast tuning phases can compensate for suboptimal configurations!



Future Work

- Dynamic Rule Generation
 - Update the expert knowledge on the fly
 - Adapt to new scenarios
 - Share rules between different use cases
- Improving Tuning Strategies
 - Implement early stopping mechanism
 - Avoid evaluating extremely bad configurations
- Simplification of the Fuzzy System to Decision Trees
 - Use Decision Trees instead of Fuzzy Systems
 - Simplify the tuning process
 - Make the tuning process more transparent

Conclusion

- Fuzzy Logic is very promising for tuning AutoPas
- Data-Driven Rule Extraction is a powerful tool
- However:
 - Requires a lot of prior data (current training data is 1.1GB)
 - Users cannot be expected to collect this data beforehand
 - A more user-friendly approach is needed for broader adoption
- Future work could focus on simplifying the process
- Alternatively: Investigate Early Stopping, to solve the tuning overhead once and for all

Thank you for your attention!

Questions?

References I

-  Keeley Crockett, Zuhair Bandar, David Mclean, and James OShea.
On constructing a fuzzy inference framework using crisp decision trees.
Fuzzy Sets and Systems, 157(21):2809–2832, 2006.
-  Ali Mohammed, Jonas H. MÄ¼ller KorndÃ¶rfer, Ahmed Eleiemy, and Florina M. Ciorba.
Automated scheduling algorithm selection and chunk parameter calculation in openmp.
IEEE Transactions on Parallel and Distributed Systems, 33(12):4383–4394, 2022.
-  Samuel James Newcome, Fabio Alexander Gratl, Philipp Neumann, and Hans-Joachim Bungartz.
Towards the smarter tuning of molecular dynamics simulations, Feb 2023.
Amsterdam, The Netherlands.